

Linear Regression, GAM Models, and RMSE Comparison

2025-05-05

Set Up Data

```
# Get data
wineData <- read.csv("~/Downloads/wine-quality-white-and-red.csv")

wineData$type <- as.factor(wineData$type)

# Split to training and testing data
set.seed(1)
train = sample(1:nrow(wineData), 0.7 * nrow(wineData))
train_data = wineData[train, ]
test_data = wineData[-train, ]
```

Best Subset Selection

```
regfit.best=regsubsets(quality~.,data=train_data, nvmax=12)
test.mat=model.matrix(quality~.,data=test_data) # create an X matrix of test data
val.errors=rep(NA,19)
for(i in 1:12){
  coefi=coef(regfit.best,id=i)
  pred=test.mat[,names(coefi)]%*%coefi
  val.errors[i] <- mean((test_data$quality - pred)^2)
}

summary(regfit.best)
```

```
## Subset selection object
## Call: regsubsets.formula(quality ~ ., data = train_data, nvmax = 12)
## 12 Variables (and intercept)
##               Forced in Forced out
## typewhite      FALSE      FALSE
## fixed.acidity   FALSE      FALSE
## volatile.acidity FALSE      FALSE
## citric.acid     FALSE      FALSE
## residual.sugar  FALSE      FALSE
## chlorides       FALSE      FALSE
## free.sulfur.dioxide FALSE      FALSE
## total.sulfur.dioxide FALSE      FALSE
## density         FALSE      FALSE
## pH              FALSE      FALSE
```

```

## sulphates          FALSE      FALSE
## alcohol            FALSE      FALSE
## 1 subsets of each size up to 12
## Selection Algorithm: exhaustive
##      typewhite fixed.acidity volatile.acidity citric.acid residual.sugar
## 1 ( 1 ) " "      " "          " "              " "          " "
## 2 ( 1 ) " "      " "          "*"              " "          " "
## 3 ( 1 ) " "      " "          "*"              " "          " "
## 4 ( 1 ) " "      " "          "*"              " "          "*"
## 5 ( 1 ) " "      " "          "*"              " "          "*"
## 6 ( 1 ) " "      " "          "*"              " "          "*"
## 7 ( 1 ) " "      " "          "*"              " "          "*"
## 8 ( 1 ) "*"      " "          "*"              " "          "*"
## 9 ( 1 ) "*"      "*"          "*"              " "          "*"
## 10 ( 1 ) "*"      "*"          "*"              " "          "*"
## 11 ( 1 ) "*"      "*"          "*"              " "          "*"
## 12 ( 1 ) "*"      "*"          "*"              "*"          "*"
##      chlorides free.sulfur.dioxide total.sulfur.dioxide density pH
## 1 ( 1 ) " "      " "          " "              " "          " "
## 2 ( 1 ) " "      " "          " "              " "          " "
## 3 ( 1 ) " "      " "          " "              " "          " "
## 4 ( 1 ) " "      " "          " "              " "          " "
## 5 ( 1 ) " "      "*"          "*"              " "          " "
## 6 ( 1 ) " "      "*"          "*"              " "          " "
## 7 ( 1 ) "*"      "*"          "*"              " "          " "
## 8 ( 1 ) " "      "*"          "*"              "*"          " "
## 9 ( 1 ) " "      "*"          " "              "*"          "*"
## 10 ( 1 ) " "      "*"          "*"              "*"          "*"
## 11 ( 1 ) "*"      "*"          "*"              "*"          "*"
## 12 ( 1 ) "*"      "*"          "*"              "*"          "*"
##      sulphates alcohol
## 1 ( 1 ) " "      "*"
## 2 ( 1 ) " "      "*"
## 3 ( 1 ) "*"      "*"
## 4 ( 1 ) "*"      "*"
## 5 ( 1 ) " "      "*"
## 6 ( 1 ) "*"      "*"
## 7 ( 1 ) "*"      "*"
## 8 ( 1 ) "*"      "*"
## 9 ( 1 ) "*"      "*"
## 10 ( 1 ) "*"      "*"
## 11 ( 1 ) "*"      "*"
## 12 ( 1 ) "*"      "*"

```

```
which.min(val.errors)
```

```
## [1] 12
```

```
coef(regfit.best,12)
```

```

##      (Intercept)          typewhite      fixed.acidity
##      1.190237e+02      -3.734753e-01      9.607001e-02
##      volatile.acidity      citric.acid      residual.sugar

```

```
##      -1.381747e+00      -1.535991e-02      6.575127e-02
##      chlorides  free.sulfur.dioxide total.sulfur.dioxide
##      -8.970624e-01      5.661084e-03      -1.625646e-03
##      density      pH      sulphates
##      -1.184895e+02      5.766522e-01      6.741088e-01
##      alcohol
##      2.107583e-01
```

Explore Predictors for GAM

```
library(ggplot2)

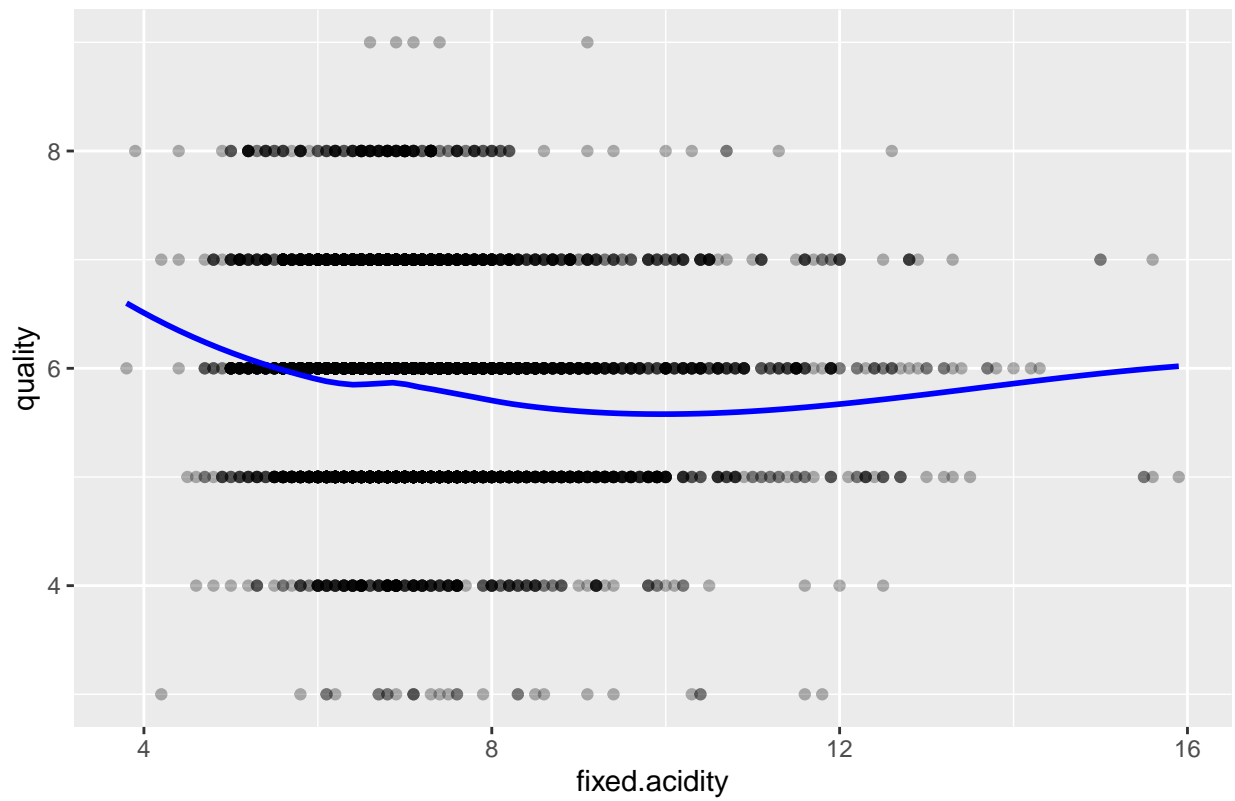
# Continuous variables
predictors <- c("fixed.acidity", "volatile.acidity", "citric.acid", "residual.sugar",
               "chlorides", "free.sulfur.dioxide", "total.sulfur.dioxide",
               "density", "pH", "sulphates", "alcohol")

# Plot each predictor against quality
for (var in predictors) {
  p <- ggplot(wineData, aes_string(x = var, y = "quality")) +
    geom_point(alpha = 0.3) +
    geom_smooth(method = "loess", se = FALSE, color = "blue") +
    labs(title = paste("Quality vs", var))
  print(p)
}
```

```
## Warning: 'aes_string()' was deprecated in ggplot2 3.0.0.
## i Please use tidy evaluation idioms with 'aes()'.
## i See also 'vignette("ggplot2-in-packages")' for more information.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was
## generated.


## 'geom_smooth()' using formula = 'y ~ x'
```

Quality vs fixed.acidity

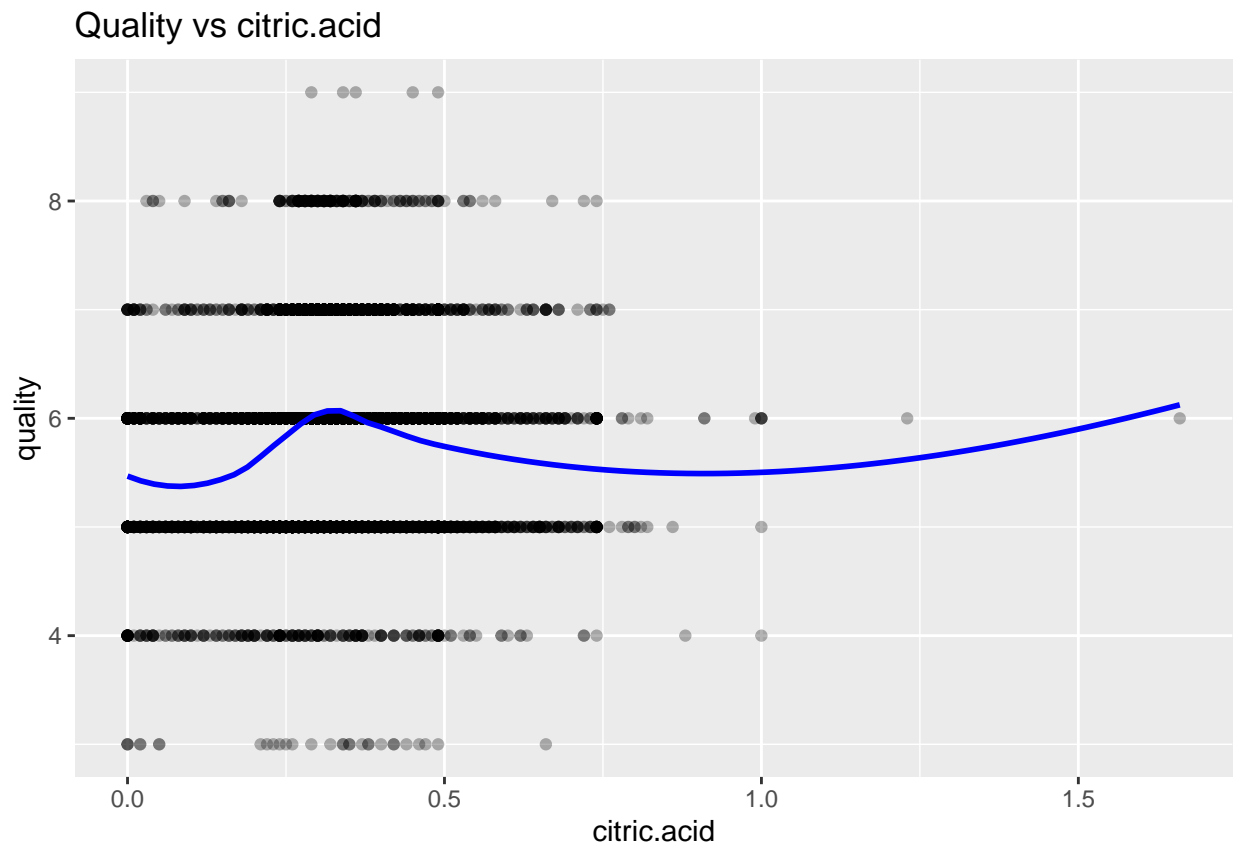


```
## 'geom_smooth()' using formula = 'y ~ x'
```

Quality vs volatile.acidity




```
## 'geom_smooth()' using formula = 'y ~ x'
```




```
## 'geom_smooth()' using formula = 'y ~ x'
```

Quality vs residual.sugar



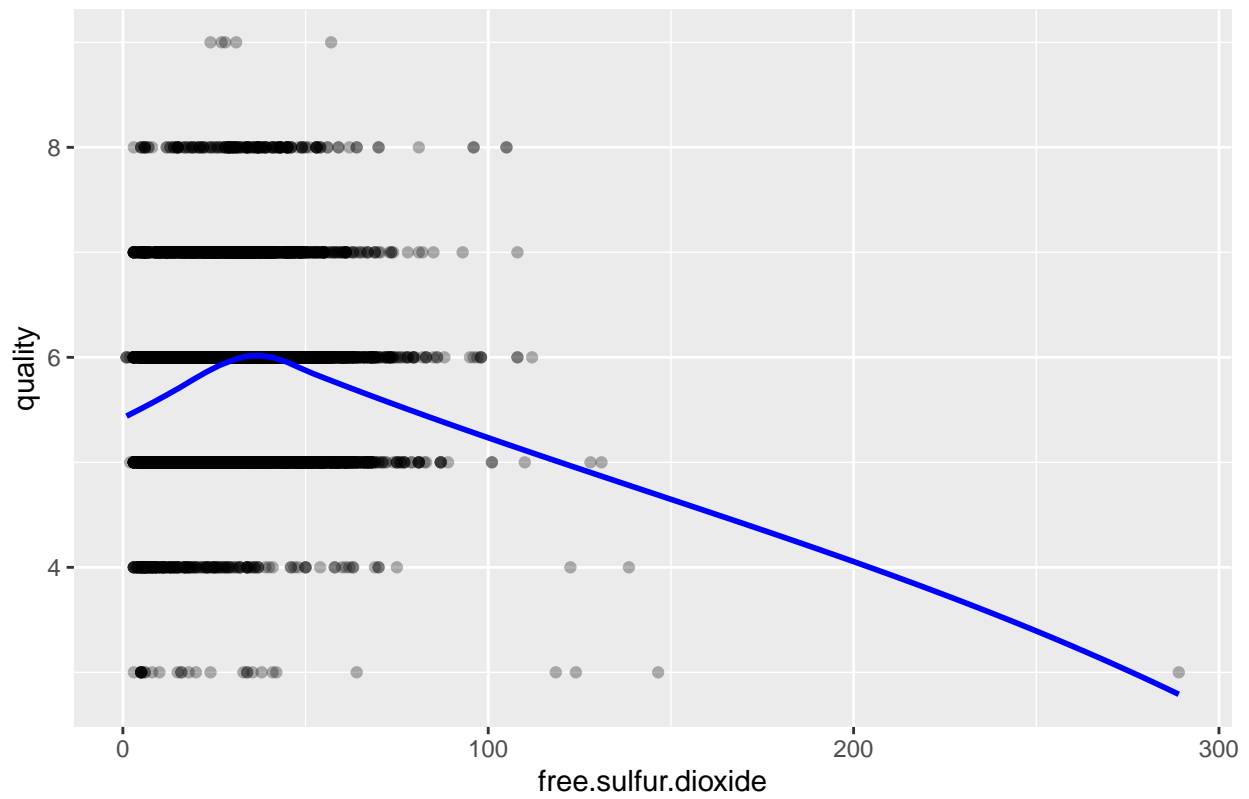
```
## 'geom_smooth()' using formula = 'y ~ x'
```

Quality vs chlorides




```
## 'geom_smooth()' using formula = 'y ~ x'
```


Quality vs free.sulfur.dioxide




```
## 'geom_smooth()' using formula = 'y ~ x'
```

Quality vs total.sulfur.dioxide




```
## 'geom_smooth()' using formula = 'y ~ x'
```

Quality vs density




```
## 'geom_smooth()' using formula = 'y ~ x'
```

Quality vs pH




```
## 'geom_smooth()' using formula = 'y ~ x'
```

Quality vs sulphates




```
## 'geom_smooth()' using formula = 'y ~ x'
```

Quality vs alcohol



```
# Categorical: type  
ggplot(wineData, aes(x = type, y = quality)) +  
  geom_boxplot() +  
  labs(title = "Quality vs Type")
```



GAM Model 1 Using Train/Test Split

```
library(gam)
```

```
## Loading required package: splines
```

```
## Loading required package: foreach
```

```
## Loaded gam 1.22-5
```

```
##
```

```
## Attaching package: 'gam'
```

```
## The following objects are masked from 'package:mgcv':
```

```
##
```

```
## gam, gam.control, gam.fit, s
```

```
# Example GAM model using smoothing for a few predictors
```

```
gam_model1 <- gam(quality ~  
  s(fixed.acidity) +  
  s(volatile.acidity) +  
  s(citric.acid) +
```

```

s(chlorides) +
s(residual.sugar) +
s(free.sulfur.dioxide) +
s(total.sulfur.dioxide) +
s(density) +
s(sulphates) +
s(alcohol) + s(pH), data = train_data)

summary(gam_model1)

##
## Call: gam(formula = quality ~ s(fixed.acidity) + s(volatile.acidity) +
##      s(citric.acid) + s(chlorides) + s(residual.sugar) + s(free.sulfur.dioxide) +
##      s(total.sulfur.dioxide) + s(density) + s(sulphates) + s(alcohol) +
##      s(pH), data = train_data)
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.5824 -0.4591 -0.0195  0.4492  2.4936
##
## (Dispersion Parameter for gaussian family taken to be 0.5009)
##
##      Null Deviance: 3461.128 on 4546 degrees of freedom
## Residual Deviance: 2254.917 on 4502 degrees of freedom
## AIC: 9806.769
##
## Number of Local Scoring Iterations: NA
##
## Anova for Parametric Effects
##              Df Sum Sq Mean Sq F value    Pr(>F)
## s(fixed.acidity)      1   14.10   14.097   28.1449 1.180e-07 ***
## s(volatile.acidity)    1  156.02  156.022  311.5014 < 2.2e-16 ***
## s(citric.acid)        1    0.01    0.011    0.0223  0.8814
## s(chlorides)          1   34.46   34.465   68.8101 < 2.2e-16 ***
## s(residual.sugar)     1   64.10   64.102  127.9807 < 2.2e-16 ***
## s(free.sulfur.dioxide) 1    0.03    0.030    0.0597  0.8070
## s(total.sulfur.dioxide) 1  165.02  165.024  329.4739 < 2.2e-16 ***
## s(density)            1  229.84  229.839  458.8803 < 2.2e-16 ***
## s(sulphates)          1  109.76  109.758  219.1344 < 2.2e-16 ***
## s(alcohol)            1  150.95  150.947  301.3690 < 2.2e-16 ***
## s(pH)                 1    8.22    8.223   16.4167 5.169e-05 ***
## Residuals           4502 2254.92    0.501
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##              Npar Df Npar F      Pr(F)
## (Intercept)
## s(fixed.acidity)      3  3.705  0.011185 *
## s(volatile.acidity)    3 12.581 3.450e-08 ***
## s(citric.acid)        3  7.803 3.414e-05 ***
## s(chlorides)          3  2.763 0.040500 *
## s(residual.sugar)     3 12.193 6.056e-08 ***
## s(free.sulfur.dioxide) 3 53.123 < 2.2e-16 ***

```



```
## s(total.sulfur.dioxide)      3 20.407 3.939e-13 ***
## s(density)                   3  4.033  0.007099 **
## s(sulphates)                 3  2.151  0.091701 .
## s(alcohol)                   3 13.383 1.077e-08 ***
## s(pH)                        3  8.208 1.909e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# Predict
gam_pred <- predict(gam_model1, newdata = test_data)

# RMSE
mse <- mean((gam_pred - test_data$quality)^2)
rmse <- sqrt(mse)
cat("RMSE of GAM1: ", rmse)
```

```
## RMSE of GAM1:  0.7220457
```

GAM Model 2 Using Train/Test Split

```
library(gam)

# Example GAM model using smoothing for a few predictors
gam_model2 <- gam(quality ~
  fixed.acidity +
  volatile.acidity +
  s(citric.acid) +
  residual.sugar +
  s(chlorides) +
  free.sulfur.dioxide +
  total.sulfur.dioxide +
  density +
  s(sulphates) +
  s(alcohol) + s(pH),
  data = train_data)

summary(gam_model2)
```

```
##
## Call: gam(formula = quality ~ fixed.acidity + volatile.acidity + s(citric.acid) +
##      residual.sugar + s(chlorides) + free.sulfur.dioxide + total.sulfur.dioxide +
##      density + s(sulphates) + s(alcohol) + s(pH), data = train_data)
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.60483 -0.47131 -0.02554  0.44413  2.59545
##
## (Dispersion Parameter for gaussian family taken to be 0.5276)
##
##      Null Deviance: 3461.128 on 4546 degrees of freedom
## Residual Deviance: 2384.858 on 4520 degrees of freedom
## AIC: 10025.52
```

```
##
## Number of Local Scoring Iterations: NA
##
## Anova for Parametric Effects
##           Df  Sum Sq Mean Sq  F value    Pr(>F)
## fixed.acidity      1    19.73   19.727   37.3892 1.049e-09 ***
## volatile.acidity    1   190.77  190.767  361.5598 < 2.2e-16 ***
## s(citric.acid)      1     0.46    0.457    0.8669   0.3519
## residual.sugar      1    46.43   46.428   87.9941 < 2.2e-16 ***
## s(chlorides)        1    42.38   42.383   80.3286 < 2.2e-16 ***
## free.sulfur.dioxide  1     0.19    0.192    0.3641   0.5462
## total.sulfur.dioxide  1   122.54  122.538  232.2448 < 2.2e-16 ***
## density             1   236.93  236.926  449.0442 < 2.2e-16 ***
## s(sulphates)        1   115.27  115.271  218.4721 < 2.2e-16 ***
## s(alcohol)          1   175.43  175.427  332.4843 < 2.2e-16 ***
## s(pH)               1    12.39   12.393   23.4886 1.299e-06 ***
## Residuals          4520 2384.86    0.528
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##           Npar Df  Npar F      Pr(F)
## (Intercept)
## fixed.acidity
## volatile.acidity
## s(citric.acid)          3  6.3365 0.0002771 ***
## residual.sugar
## s(chlorides)           3  2.6740 0.0456893 *
## free.sulfur.dioxide
## total.sulfur.dioxide
## density
## s(sulphates)           3  2.0634 0.1028790
## s(alcohol)             3 13.6932 6.868e-09 ***
## s(pH)                  3  9.2017 4.567e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# Predict
gam_pred2 <- predict(gam_model2, newdata = test_data)

# RMSE
mse <- mean((gam_pred2 - test_data$quality)^2)
rmse <- sqrt(mse)
cat("RMSE of GAM2: ", rmse)
```

```
## RMSE of GAM2: 0.7348539
```

Comparing RMSE using 10-Fold Cross Validation for Linear Regression, GAM1, GAM2, KNN, and Lasso

```

# Get data
wineData <- read.csv("~/Downloads/wine-quality-white-and-red.csv")

wineData$type <- as.factor(wineData$type)

# Split to training and testing data
set.seed(1)
train = sample(1:nrow(wineData), 0.7 * nrow(wineData))
train_data = wineData[train, ]
test_data = wineData[-train, ]

# Normalize function for KNN
normalize <- function(x) {
  return((x - mean(x)) / sd(x))
}

wine <- wineData

# Normalize numeric columns except quality
all_columns <- names(wine)
columns_to_normalize <- all_columns[all_columns != "quality" & sapply(wine, is.numeric)]
wine[columns_to_normalize] <- lapply(wine[columns_to_normalize], normalize)

# Convert factor
wine$type <- as.numeric(factor(wine$type))

# Setup cross-validation
k <- 10
folds <- sample(1:k, nrow(wine), replace = TRUE)

# Empty vectors to store RMSE
rmse_linear <- rmse_gam1 <- rmse_gam2 <- rmse_knn <- rmse_lasso <- numeric(k)

for (i in 1:k) {
  ktrain_data <- wine[folds != i, ]
  ktest_data <- wine[folds == i, ]

  # Linear Regression
  lm_model <- lm(quality ~ ., data = ktrain_data)
  linear_preds <- predict(lm_model, newdata = ktest_data)
  rmse_linear[i] <- sqrt(mean((linear_preds - ktest_data$quality)^2))

  # GAM Model 1 (smooth all variables)
  gam_model1 <- gam(quality ~
    s(fixed.acidity) +
    s(volatile.acidity) +
    s(citric.acid) +
    s(chlorides) +
    s(residual.sugar) +
    s(free.sulfur.dioxide) +
    s(total.sulfur.dioxide) +
    s(density) +

```

```

        s(sulphates) +
        s(alcohol) + s(pH), data = ktrain_data, select=TRUE)
gam_preds1 <- predict(gam_model1, newdata = ktest_data)
rmse_gam1[i] <- sqrt(mean((gam_preds1 - ktest_data$quality)^2))

# GAM Model 2 (smooth most variables)
gam_model2 <- gam(quality ~
  fixed.acidity +
  volatile.acidity +
  s(citric.acid) +
  residual.sugar +
  s(chlorides) +
  free.sulfur.dioxide +
  total.sulfur.dioxide +
  density +
  s(sulphates) +
  s(alcohol) + s(pH),
  data = ktrain_data)
gam_preds2 <- predict(gam_model2, newdata = ktest_data)
rmse_gam2[i] <- sqrt(mean((gam_preds2 - ktest_data$quality)^2))

# KNN
train_knn <- ktrain_data
test_knn <- ktest_data
train_knn$quality <- as.factor(train_knn$quality)
test_knn$quality <- as.factor(test_knn$quality)

training_X <- dplyr::select(train_knn, -quality)
testing_X <- dplyr::select(test_knn, -quality)

knn_preds <- knn.reg(train = training_X, test = testing_X, y = as.numeric(train_knn$quality), k = 10)
rmse_knn[i] <- sqrt(mean((knn_preds - as.numeric(test_knn$quality))^2))

# Lasso
x_train <- as.matrix(dplyr::select(ktrain_data, -quality))
y_train <- ktrain_data$quality
x_test <- as.matrix(dplyr::select(ktest_data, -quality))

lasso_cv <- cv.glmnet(x_train, y_train, alpha = 1)
best_lambda <- lasso_cv$lambda.min
lasso_preds <- predict(lasso_cv, s = best_lambda, newx = x_test)
rmse_lasso[i] <- sqrt(mean((lasso_preds - ktest_data$quality)^2))
}

# Average RMSE across folds
results <- data.frame(
  Model = c("Linear", "GAM1", "GAM2", "KNN", "Lasso"),
  RMSE = c(mean(rmse_linear), mean(rmse_gam1), mean(rmse_gam2), mean(rmse_knn), mean(rmse_lasso))
)

print(results)

```

##	Model	RMSE
## 1	Linear	0.7336196
## 2	GAM1	0.7149347
## 3	GAM2	0.7291629
## 4	KNN	0.6903177
## 5	Lasso	0.7336021